



Content Course Tools ▾ Student Support Contact Support

Assignments › Practical Lab 1 - Streaming Data for Predictive Maintenance with Linear Regression-Based Alerts

# Practical Lab 1 - Streaming Data for Predictive Maintenance with Linear Regression-Based Alerts

▼ Hide Assignment Information

Instructions

## Practical Lab 1: Streaming Data for Predictive Maintenance with Linear Regression-Based Alerts

### Context

In the **Data Stream Visualization Workshop**, you learned how to stream, store, and visualize industrial current data.

This lab extends that work: you will now implement **regression-based anomaly detection** to generate **alerts and errors** when currents deviate significantly from their expected values.

This task simulates a **Predictive Maintenance** scenario, where early alerts and errors can flag potential failures before they occur.

---

### Learning Objectives

By the end of this assignment, you will be able to:

- Extend a streaming pipeline with **machine learning models**.
- Apply **linear regression** to detect unusual consumption trends.
- Analyze regression residuals and discover meaningful thresholds for anomaly detection.
- Implement an **alerts/errors module** in a streaming context.

# Deliverables

Your GitHub repository must contain:

## 1. README.md

- Professional and well-structured.
- Project summary.
- Clear setup instructions.
- Explanation of regression/alert rules.
- Screenshots or plots of results.

## 2. requirements.txt

- All dependencies with versions.

## 3. Data folder

- Relevant CSV files used for training (e.g., RMBR4-2\_export\_test.csv).

## 4. Codebase

Must include scripts and/or Jupyter Notebooks that:

- Connect to a **cloud-based database** (e.g., Neon.tech PostgreSQL). Use the database to pull the information in it to the application and train the Linear Regression model.
- Ingest and query streaming data. You can read the testing data from the synthetically generated data in a Dataframe or from a CSV file.
- Run regression models (Time → Axes #1–#8).
- Implement alerts/errors based on thresholds you discover.

## 🔍 Discovering Thresholds

You are not given fixed thresholds. Instead, you must **discover them from the data**.

To do this:

### 1. Develop regression models:

- For each axis (#1–#8), fit (**train**) a univariate linear regression (Time → Axis values) and produce a model able to make predictions when fed **testing** data (read below)
- Record slope and intercept.
- Plot scatter data with regression lines.

## 2. Analyze residuals:

- Compute the difference between observed values and the regression prediction.
- Plot distributions (scatter, line, or boxplots of residuals).
- Look for outliers and patterns.

## 3. Define thresholds:

- Choose **MinC**: the minimum current deviation (kWh above regression line) that should trigger an **Alert** if sustained.
- Choose **MaxC**: the maximum current deviation (kWh above the regression line) that should trigger an **Error** if sustained.
- Choose **T**: the minimum continuous time (in seconds) that the deviation must persist.

## 4. Implement alert/error rules:

- **Alert**:  $\geq \text{MinC}$  kWh above regression line for  $\geq T$  seconds continuously
- **Error**:  $\geq \text{MaxC}$  kWh above regression line for  $\geq T$  seconds continuously

## 5. Produce testing data *synthetically*

- Use the **training metadata** to synthetically generate **testing** data to run predictions using the linear regression model.

## 6. Visualize alerts/errors:

- Overlay Alert/Error markers on your regression plots.
- Annotate each with its duration.

## 7. Log results:

- Store Alert and Error events in a structured CSV or database table.

👉 **Hint:** To produce the testing data:

- Upload the training data to an LLM
- Write a prompt to request that the LLM generate a new file with the same metadata, and a similar mean/std. deviation.
- Normalize the data set, if necessary.
  - Min-Max Normalization Algorithm
- Standardize the data set, if necessary
  - Use Z-scores to find proportions

 **Hint:** Follow the process we explored in the workshop extension:

- Fit regression models across all axes,
- Compare residuals and outlier counts,
- Test different thresholds and time windows until alerts/errors appear,
- Interpret your results in the context of **Predictive Maintenance**.

You will be graded not only on whether your code works, but also on **how well you justify your chosen MinC, MaxC, and T values** using evidence from your analysis.

---

## Submission Instructions

- Push your completed project to your GitHub repository.
  - Submit a PDF with your name, student IDs, and the URL to your repository via the course website.
  - Your GitHub repository must include all deliverables listed above.
- 

## Marking Rubric (10 points total)

	Category	Criteria	Points
Project Setup		Repo includes README.md, requirements.txt, and data/. README files that are professional and clear. The IPYNB file has the output of the last test run before submitting the project.	1
Database Integration		Code connects to Neon.tech (or equivalent PostgreSQL) and ingests/queries streaming data.	1.5
Streaming Simulation		Script simulates CSV → DB time-based flow using 1 synthetic test data that has been properly	

Category	Criteria	Points
Regression Models & Residual Analysis	normalized and standardized with respect to the training data from the relational database.	
Threshold Discovery & Justification	Fits univariate regressions (Time → Axes #1–#8), plots regression lines, and analyzes residuals.	2
Alerts & Errors Implementation	Student justifies chosen <b>MinC</b> , <b>MaxC</b> , <b>T</b> with evidence from analysis (plots, residuals, predictive maintenance context).	2
Visualization/Dashboard	Correctly implements detection logic with chosen thresholds. Logs events.	2
	Shows regression with alert/error annotations. Clear plots.	0.5

**Total: 10 points**

---

This assignment is a **follow-up** to the Data Streaming Workshop:

- You reuse the pipeline they built,
  - Extend it with regression + thresholds,
  - And must **think critically** about thresholds instead of copying fixed values.
- 

Submissions past the due date:

- 10% deduction after 11:00 AM pm on the due date.
  - 25% deduction after 11:59 pm on the due date.
  - No credit after 11:59 pm on the day after the due date.
- 

Academic integrity:

- **The instructor may summon students suspected of academic dishonesty for a code review to verify their understanding and ability to reproduce any stage of the development process. The instructor reserves the right to deduct marks, either partially or entirely.**

Due on Feb 6, 2026 11:59 PM

Available on Jan 9, 2026 12:01 AM. Access restricted before availability starts.

Available until Feb 8, 2026 11:59 PM. Access restricted after availability ends.

# Submit Assignment

Allowed File Extensions

pdf

Files to submit

(0) file(s) to submit

After uploading, you must click Submit to complete the submission.

[Add a File](#)

[Record Audio](#)

[Record Video](#)

Comments